

## TTIC 31230 Fundamentals of Deep Learning, winter 2019

### Backpropagation Problems

**Problem 0:** The KL-divergence between two discrete distributions is defined by

$$KL(P, Q) = E_{x \sim P} \ln \frac{P(x)}{Q(x)}$$

We will show later in the class that  $KL(P, Q) \geq 0$  for any  $P$  and  $Q$ .  
The Cross Entropy  $H(P, Q)$  is defined by

$$H(P, Q) = E_{x \sim P} - \ln Q(x)$$

When  $P$  is a population distribution and  $Q$  is a model distribution this is the cross entropy loss.

The entropy of a distribution is defined by

$$H(P) = E_{x \sim P} - \ln P(x)$$

Show the following.

$$H(P, Q) \geq H(P)$$

$$H(P, P) = H(P)$$

Explain why these facts are important when training a model to minimize cross entropy loss.

**Problem 1: Backpropagation through softmax.** Consider the following softmax.

$$\begin{aligned} Z[b] &= \sum_j \exp(s[b, j]) \\ p[b, j] &= \exp(s[b, j]) / Z[b] \end{aligned}$$

An alternative way to compute this is to initialize the tensors  $Z$  and  $p$  to zero and then execute the following loops.

for  $b, j$      $Z[b] += \exp(s[b, j])$

for  $b, j$      $p[b, j] += \exp(s[b, j]) / Z[b]$

Each individual  $+=$  operation inside the loops can be treated independently in backpropagation.

(a) Give a back-propagation loop over  $+=$  updates based on the second loop for adding to  $s.grad$  using  $p.grad$  (and using the forward-computed tensors  $Z$  and  $s$ ).

(b) Give a back-propagation loop over += updates based on the second equation for adding to  $Z.\text{grad}$  using  $p.\text{grad}$  (and using the forward-computed tensors  $s$  and  $Z$ ).

(c) Give a back-propagation loop over += updates based on the first equation for adding to  $s.\text{grad}$  using  $Z.\text{grad}$  (and using the forward-computed tensor  $s$ ).

**Problem 2: Optimizing Backpropagation through softmax.** Show that the addition to  $s.\text{grad}$  shown in problem 1 can be computed using the following more efficient updates.

for  $b, j$   $e[b] -= p[b, j]p.\text{grad}[b, j]$

for  $b, j$   $s.\text{grad}[b, j] += p[b, j](p.\text{grad}[b, j] + e[b])$

This formula shows how hand-written back-propagation methods for “layers” such as softmax can be more efficient than compiler-generated back-propagation code. While optimizing compilers can of course be written, one must keep in mind the trade-off between the abstraction level of the programming language and the efficiency of the generated code.

**Problem 3. Backpropagation through batch normalization.** Consider the following set of += statements defining batch normalization where all computed tensors are initialized to zero.

For  $b, j$   $\mu[j] += \frac{1}{B} x[b, j]$

For  $b, j$   $s[j] += \frac{1}{B-1} (x[b, j] - \mu[j])^2$

For  $b, j$   $x'[b, j] += \frac{x[b, j] - \mu[j]}{\sqrt{s[j]}}$

Give backpropagation += (or -=) loops for computing  $x.\text{grad}[b, j]$ ,  $\mu.\text{grad}[j]$ , and  $s.\text{grad}[j]$  from  $x'.\text{grad}[b, j]$ . The loops should be given in the order they are to be executed.

**Problem 4. Backpropagation through a UGRNN.** Equations defining a UGRNN are given below.

$$\tilde{R}_t[b, j] = \left( \sum_i W^{h,R}[j, i] h_{t-1}[b, i] \right) + \left( \sum_k W^{x,R}[j, k] x_t[b, k] \right) - B^R[j]$$

$$R_t[b, j] = \tanh(\tilde{R}_t[b, j])$$

$$\tilde{G}_t[b, j] = \left( \sum_i W^{h,G}[j, i] h_{t-1}[b, i] \right) + \left( \sum_k W^{x,G}[j, k] x_t[b, k] \right) - B^G[j]$$

$$G_t[b, j] = \sigma(\tilde{G}_t[b, j])$$

$$h_t[b, j] = G_t[b, j] h_{t-1}[b, j] + (1 - G_t[b, j]) R_t[b, j]$$

- (a) Rewrite the first equation defining  $\tilde{R}_t$  using += loops instead of summations assuming that all computed tensors are initialized to zero.
- (b) Give += loops for the backward computation for your solution to part (a) using the convention that parameter gradients are averaged over the batch and where the batch size is  $B$ .

**Problem 5. Writing framework code.** Consider a function  $c : R^d \times R^s \rightarrow R^s$ , in other words a function that takes a vector of dimension  $d$  and a vector of dimension  $s$  and yields a vector of dimension  $s$ . Given a sequence of vectors  $x_0, x_2, \dots, x_T$  with  $x_t \in R^d$  we can define a sequence of vectors  $h_0, h_1, \dots, h_T$  by the equations

$$\begin{aligned} h_0 &= c(x_0, 0) \\ h_t &= c(x_t, h_{t-1}) \text{ for } 1 \leq t \leq T \end{aligned}$$

When the function  $c$  is defined by a neural network the resulting network mapping  $x_1, \dots, x_T$  to  $h_0, \dots, h_T$  is called a recurrent neural network (RNN).

**a.** In the educational framework EDF we work with objects where each object has a value attribute and a gradient attribute each of which have tensor values where the value tensor and the gradient tensor are the same shape. Each object is assigned a value in a forward pass and assigned a gradient in a backward pass. Suppose that we are given an EDF procedure `CELL` which takes as arguments a parameter object `Phi` and two EDF objects `X` and `H` where the value attribute of the object `X` is a  $d$ -dimensional vector and the value attribute of the object `H` is an  $s$ -dimensional vector. A call to the procedure `CELL(Phi, X, H)` returns an EDF object whose value attribute is computed in a forward pass in some possibly complex way from the value attributes of `Phi`, `X` and `H`. Given a sequence `X[]` of EDF objects whose value attributes are  $d$ -dimensional vectors, and an EDF object `ZERO` representing the constant  $s$ -dimensional zero vector, write a procedure for constructing the sequence of EDF objects representing  $h_1, h_2, \dots, h_T$  as defined by the above RNN equations. Your solution can be in Python or informal high level pseudo code.

**b.** Deep learning systems generally make extensive use of parallel computation for training. How does the parallel running time of an RNN computation graph scale with the length  $T$ ?