Problems For Language Modeling, Translation and Attention.

Problem 1. Blank Language modeling. This problem considers "blank language modeling" which is used in BERT. For blank language modelng we draw a sentence $w_1, \ldots, w_T$ from a corpose and blank out a word at random and ask the system to predict the blanked word. The cross-entropy loss for blank language modeling can be written as

$$\Phi^* = \underset{\Phi}{\mathrm{argmin}} \ E_{w_1,\ldots,w_T \sim \mathrm{Train}, t \sim \{1,\ldots,T\}} \ \ -\ln P_\Phi(w_t | w_1, \ldots, w_{t-1}, w_{t+1}, \ldots, w_T)$$

Consider a bidirectional RNN run on a sequence of words $w_1, \ldots, w_T$ such that for each time t we have a forward hidden state $\vec{h}[t, J]$ computed from $w_1, \ldots, w_t$ and a backward hidden state $\overleftarrow{h}[t, J]$ computed from $w_T, \ w_{T-1}, \ldots w_t$. Also assume that each word w has an associated word vector $e[w, J]$. Give a definition of $P(w_t \mid w_1, \ldots w_{t-1}, \ w_{t+1}, \ldots, w_T)$ as a function of the vectors $\vec{h}[t-1, J]$ and $\overleftarrow{h}[t+1, J]$ and the word vectors $e[W, I]$. You can assume that $\vec{h}[T, J]$ and $\overleftarrow{[h]}[T, J]$ have the same shape (same dimensions) but do not make any assumptions about the dimension of the word vectors $e[W, I]$. You can assume whatever tensor parameters you want.

Problem 2. Image Captioning as Translation with Attention. In machine translation with attention the translation is generated from an autoregressive language model for the target langauge translation given the source language sentence. The source sentence is converted to an initial hidden vector $h[0, J]$ for the decoding (usually the final hidden vector of a right-to-left RNN run on the input), plus the sequence $M[T_{\mathrm{in}}, J]$ of hidden vectors computed by the RNN on the source sentence where T is the length of the source sentence. We then define an autoregressive conditional language model

$$P_\Phi(w_1, \ldots, w_{T_{\mathrm{out}}} \mid h[0, J], \ M[T_{\mathrm{in}}, J])$$

An autoregressive conditional language model with attention can be defined by

$$P(w_t \mid w_0, \cdots, w_{t-1}) \ = \ \underset{w_t}{\mathrm{softmax}} \ e[w_t, I] W^{\mathrm{auto}}[I, J] h[t-1, J]$$

$$\alpha[t_{\mathrm{in}}] \ = \ \underset{t_{\mathrm{in}}}{\mathrm{softmax}} \ h[t-1, J_1] W^{\mathrm{key}}[J_1, J_2] M[t_{\mathrm{in}}, J_2]$$

$$V[J] \ = \ \sum_{t_{\mathrm{in}}} \alpha[t_{\mathrm{in}}] M[t_{\mathrm{in}}, J]$$

$$h[t, J] \ = \ \mathrm{CELL}_\Phi(h[t-1, J], \ V[J], \ e[w_t, I])$$

Here CELL is some function taking (objects for) two vectors of dimensions J and one vector of dimension I and returning (an object for) a vector of dimension J.

Rewrite these equations for image captioning where instead of $M[t_{\text{in}}, J]$ we are given an image feature tensor $M[x, y, K]$

Problem 3. Language CNNs. This problem is on CNNs for sentences. We consider a model with parameters

$$\Phi = (e[w, i], W_1[\Delta t, i, i'], B_1[i], \ldots, W_L[\Delta t, i, i'], B_L[i])$$

The matrix e is the word embedding matrix where $e[w, I]$ is the vector embedding of word w.

(a) Give an equation for the convolution layer $L_0[b, t, i]$ as a function of the word embeddings and the input sentence $w_1, \ldots, w_T$.

(b) Give an equation for $L_{\ell+1}[b, t, i]$ as a function of $L_\ell[b, t, i]$ and the parameters $W_{\ell+1}[\Delta t, i', i]$ and $B_{\ell+1}[i]$ and where $L_{\ell+1}$ is computed stride 2.

(c) Assuming all computations can be done in parallel as soon the inputs have been computed, what is the **parallel** order of run time for this convolutional model as a function of the input length T and the number of layers L (assume all parameter tensors of size $O(1)$). Compare this with the parallel run time of an RNN.

Problem 4. A self-attention layer in the transformer takes a sequence of vectors $h_{\text{in}}[T, J]$ and computes a sequence of vectors $h_{\text{out}}[T, J]$ using the following equations where k ranges over "heads". Heads are intended to allow for different relationship between words such as "coreference" or "subject of" for a verb. But the actual meaning emerges during training and is typically difficult or impossible to interpret. In the following equations we typically hve $U < J$ and we require $I = J/K$ so that the concatenation of K vectors of dimension I is a vector of dimension J.

$$
\begin{aligned}
\text{Query}[k, t, U] &= W^Q[k, U, J] h_{\text{in}}[t, J] \\[2mm]
\text{Key}[k, t, U] &= W^K[k, U, J] h_{\text{in}}[t, J] \\[2mm]
\alpha[k, t_1, t_2] &= \underset{t_2}{\text{softmax}} \; \text{Query}[k, t_1, U] \text{Key}[k, t_2, U] \\[2mm]
\text{Value}[k, t, I] &= W^V[k, I, J] h_{\text{in}}[t, J] \\[2mm]
\text{Out}[k, t, I] &= \sum_{t'} \alpha[k, t, t'] \text{Value}[k, t', I] \\[2mm]
h_{\text{out}}[t, J] &= \text{Out}[1, t, I]; \cdots ; \text{Out}[K, t, I]
\end{aligned}
$$

A summation over N terms can be done in parallel in $O(\log N)$ time.

(a) For a given head k and position t_1 what is the parallel running time of the above softmax operation, as a function of T and U where we first compute the scores to be used in the softmax and then compute the normalizing constant Z.

(b) What is the order of running time of the self-attention layer as a function of T, J and K (we have I and U are both less than J.)

Problem 5. Just as CNNs can be done in two dimensions for vision and in one dimension for language, the Transformer can be done in two dimensions for vision — the so-called spatial transformer.

(a) Rewrite the equations from problem 1 so that the time index t is replaced by spatial dimensions x and y.

(b) Assuming that summations take logarithmic parallel time, give the parallel order of run time for the spatial self-attention layer as a function of X, Y, J and K (we have that I and U are both less than J).

Problem 6. The self-attention in the transformer is computed by the following equations.

$$\text{Query}_{\ell+1}[k, t, i] \quad = \quad W^Q_{\ell+1}[k, i, J] L_\ell[t, J]$$

$$\text{Key}_{\ell+1}[k, t, i] \quad = \quad W^K_{\ell+1}[k, i, J] L_\ell[t, J]$$

$$\alpha_{\ell+1}[k, t_1, t_2] \quad = \quad \operatorname*{softmax}_{t_2} \left[\frac{1}{\sqrt{I}} \text{Query}_{\ell+1}[k, t_1, I] \text{Key}_{\ell+1}[k, t_2, I] \right]$$

Notice that here the shape of W^Q and W^K are both $[K, I, J]$. We typically have $I < J$ which makes the inner product in the last line an inner product of lower dimensional vectors.

(a) 20 pts Give an equation computing a tensor $\tilde{W}^Q[K, J, J]$ computed from W^Q and W^K such that the attention $\alpha(k, t_1, t_2)$ can be written as

$$\alpha_{\ell+1}(k, t_1, t_2) = \operatorname*{softmax}_{t_2} \left[L_\ell[t_1, J_1] \tilde{W}^Q[k, J_1, J_2] L_\ell[t_2, J_2] \right]$$

For a fixed k we have that $W^Q[k, I, J]$ and $W^K[k, I, J]$ are matrices. We want a matrix $\tilde{W}^Q[k, J, J]$ such that the attention can be written in matrix notation as $h_1^\top \tilde{W}^Q h_2$ where h_1 and h_2 are vectors and $\tilde{W}^Q$ is a matrix. You need write this matrix $\tilde{W}^Q$ in terms of the matrices for W^Q and W^K. But write your final answer in Einstein notation with k as the first index.

(b) 5pts Part (a) shows that we can replace the key and query matrix with a single query matrix without any loss of expressive power. If we eliminate the key matrix in this way what is the resulting number of query matrix parameters for a given layer and how does this compare to the number of key-query matrix parameters for a given layer in the original transformer version.