

TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2020

Machine Translation and Attention

Machine Translation

$$w_0, \dots, w_{T_{\text{in}}} \Rightarrow \tilde{w}_0, \dots, \tilde{w}_{T_{\text{out}}}$$

Translation is a **sequence to sequence** (seq2seq) task.

Sequence to Sequence Learning with Neural Networks, Sutskever, Vinyals and Le, NeurIPS 2014, arXiv Sept 10, 2014.

We describe a simplification of the paper.

Machine Translation

We define a model

$$P_{\Phi} (\tilde{w}_0, \dots, \tilde{w}_{T_{\text{out}}} \mid w_0, \dots, w_{T_{\text{in}}})$$

$$\begin{aligned} \Phi^* &= \operatorname{argmin}_{\Phi} E_{\text{Pop}} - \ln P_{\Phi} (\tilde{w}_0, \dots, \tilde{w}_{T_{\text{out}}} \mid w_0, \dots, w_{T_{\text{in}}}) \\ &= \operatorname{argmin}_{\Phi} E_{\langle x, y \rangle \sim \text{Pop}} - \ln P_{\Phi}(y|x) \end{aligned}$$

A Simple RNN Translation Model

The final state of a **right-to-left (backward)** RNN, $\overleftarrow{h}_{\text{in}}[0, J]$, is viewed as a “**thought vector**” representation of the input sentence.

We use the input thought vector $\overleftarrow{h}_{\text{in}}[0, J]$ as the initial hidden state a **left-to-right (forward)** RNN language model generating the output sentence.

Computing the input thought vector backward provides a good start to the forward generation of the output.

Machine Translation Decoding

We can sample a translation

$$w_t \sim P(w_t \mid \overleftarrow{h}_{\text{in}}[0, J], w_0, \dots, w_{t-1})$$

or we can do greedy decoding

$$w_t = \operatorname{argmax}_{w_t} P(w_t \mid \overleftarrow{h}_{\text{in}}[0, J], w_0, \dots, w_{t-1})$$

or we might try maximize total probability.

$$w_0, \dots, w_{T_{\text{out}}} = \operatorname{argmax}_{w_0, \dots, w_{T_{\text{out}}}} P_{\Phi} \left(w_0, \dots, w_{T_{\text{out}}} \mid \overleftarrow{h}_{\text{in}}[0, J] \right)$$

Greedy Decoding vs. Beam Search

We would like

$$W_{\text{out}}[T_{\text{out}}]^* = \operatorname{argmax}_{W_{\text{out}}[T_{\text{out}}]} P_{\Phi}(W_{\text{out}}[T_{\text{out}}] \mid W_{\text{in}}[T_{\text{in}}])$$

But a greedy algorithm may do well

$$w_t = \operatorname{argmax}_{w_t} P_{\Phi}(w_t \mid W_{\text{in}}[T_{\text{in}}], w_0, \dots, w_{t-1})$$

But these are not the same.

Example

“Those apples are good” vs. “Apples are good”

$$P_{\Phi}(\text{Apples are Good } \langle \text{eos} \rangle) > P_{\Phi}(\text{Those apples are good } \langle \text{eos} \rangle)$$

$$P_{\Phi}(\text{Those}|\varepsilon) > P_{\Phi}(\text{Apples}|\varepsilon)$$

Beam Search

At each time step we maintain a list the K best words and their associated hidden vectors.

This can be used to produce a list of k “best” decodings which can then be compared to select the most likely one.

Machine Translation with Attention

**Neural Machine Translation by Jointly Learning to
Align and Translate** Dzmitry Bahdanau, Kyunghyun Cho,
Yoshua Bengio, ICLR 2015 (arXiv Sept. 1, 2014)

We describe a simplification of the paper.

Representing Sentences by Vector Sequences

We first run a bidirectional RNN on the input sentence to get a **sequence** $\overleftrightarrow{h}_{\text{in}} [T_{\text{in}}, J]$ of hidden vectors $\overleftrightarrow{h}_{\text{in}} [t_{\text{in}}, J]$ for $1 \leq t_{\text{in}} \leq T_{\text{in}}$. We then replace

$$w_t \sim P(w_t \mid \overleftrightarrow{h}_{\text{in}}[0, J], w_0, \dots, w_{t-1})$$

by

$$w_t \sim P(w_t \mid \overleftrightarrow{h}_{\text{in}} [T_{\text{in}}, J], w_0, \dots, w_{t-1})$$

The autoregressive probability model of the output is now conditioned on a sequence of vectors rather than a single thought vector.

Machine Translation with Attention

$$\text{Autoregression :} \quad P(w_{t_{\text{out}}} \mid \overset{\leftrightarrow}{h}_{\text{in}} [T_{\text{in}}, J], w_0, \dots, w_{t_{\text{out}}-1})$$

$$\text{Autoregression :} \quad = \text{softmax}_{w_{t_{\text{out}}}} e[w_{t_{\text{out}}}, J] \vec{h}_{\text{out}}[t_{\text{out}} - 1, J]$$

$$\text{RNN :} \quad \vec{h}_{\text{out}}[0, J] = \overleftarrow{h}_{\text{in}}[0, J/2]; \vec{h}_{\text{in}}[T_{\text{in}}, J/2]$$

$$\text{RNN :} \quad \vec{h}_{\text{out}}[t_{\text{out}}, J] = \text{CELL}(\vec{h}_{\text{out}}[t_{\text{out}} - 1, J], e[w_{t_{\text{out}}}, I], \tilde{h}_{\text{in}}[t_{\text{out}}, J])$$

$$\text{Attention :} \quad \tilde{h}_{\text{in}}[t_{\text{out}}, J] = \alpha[t_{\text{out}}, T_{\text{in}}] \overset{\leftrightarrow}{h}_{\text{in}} [T_{\text{in}}, J]$$

$$\text{Attention :} \quad \alpha[t_{\text{out}}, t_{\text{in}}] = \text{softmax}_{t_{\text{in}}} e[w_{t_{\text{out}}}, J] \overset{\leftrightarrow}{h}_{\text{in}} [t_{\text{in}}, J]$$

Attention

$$\alpha[t_{\text{out}}, t_{\text{in}}] = \underset{t_{\text{in}}}{\text{softmax}} \ e[w_{t_{\text{out}}}, J] \quad \overleftrightarrow{h}_{\text{in}} [t_{\text{in}}, J]$$

$$\tilde{h}_{\text{in}}[t_{\text{out}}, J] = \alpha[t_{\text{out}}, T_{\text{in}}] \quad \overleftrightarrow{h}_{\text{in}} [T_{\text{in}}, J]$$

$\tilde{h}_{\text{in}}[t_{\text{out}}, J]$ is a convex combination of vectors $\overleftrightarrow{h}_{\text{in}} [t_{\text{in}}, J]$.

More generally, attention computes a convex combination of vectors where the combination weights are computed by a softmax of an inner product with a “query” vector (such as $e[w_{t_{\text{out}}}, J]$ above).

Attention in Image Captioning

We can treat image captioning as translating an image into a caption.


In translation with attention involves an attention over the input aligning output words with positions in the input.

For each output word we get an attention over the image positions.

Attention in Image Captioning



A woman is throwing a frisbee in a park.



A little girl sitting on a bed with a teddy bear.

Xu et al. ICML 2015

END