TTIC 31230, Fundamentals of Deep Learning

David McAllester, Winter 2021

Gaussian Models

and The Parels of Differential Entropy

Gaussian VAEs

$$\text{VAE:} \quad \Phi^*, \Psi^* = \operatorname*{argmin}_{\Phi, \Psi} \; E_{y,z} \; \ln \frac{\hat{p}_\Psi(z|y)}{p_\Phi(z)} - \ln p_\Phi(y|z)$$

All models are Gaussian densities.

$$p_\Phi(z[i]) \; \propto \; \exp((z[i] - \mu_\Phi[i])^2 / 2\sigma_\Phi^2[i])$$

$$p_\Psi(z[i]|y) \; \propto \; \exp((z[i] - \hat{z}_\Psi(y)[i])^2 / 2\sigma_\Psi^2(y)[i])$$

$$p_\Phi(y[i]|z) \; \propto \; \exp((y[i] - \hat{y}_\Phi(z)[i])^2 / 2\sigma_\Phi^2(z)[i])$$

Differential Entropy

In the case of a continuous density (as opposed to a discrete probability) we have the notion of differential entropy.

For a density $p(x)$ on a real value x we have

$$H(p) = E_{x \sim p} \left[-\ln p(x) \right]$$

$$= \int_{-\infty}^{\infty} p(x) \left(-\ln p(x) \right) dx$$

Differential Entropy can Diverge to $-\infty$

For a uniform distribution over an interval on the real line of width Δ we have

$$H = E_{x \sim p}\left[-\ln p(x)\right]$$

$$= E_{x \sim p}\left[-\ln \frac{1}{\Delta}\right]$$

$$= \ln \Delta$$

As $\Delta \to 0$ we have $H \to -\infty$

Differential Entropy can Diverge to $-\infty$

$$\begin{aligned}
H(\mathcal{N}(0, \sigma^2)) &= E_{x \sim \mathcal{N}(0,\sigma^2)} \left[-\ln \left(\frac{1}{\sqrt{\pi}\sigma} \exp \frac{-x^2}{2\sigma^2} \right) \right] \\
&= E_{x \sim \mathcal{N}(0,\sigma^2)} \left[\ln(\sqrt{\pi}\sigma) + \frac{-x^2}{2\sigma^2} \right] \\
&= (\ln \sigma) + \ln(\sqrt{\pi}) + E_x \left[\frac{x^2}{2\sigma^2} \right] \\
&= (\ln \sigma) + \ln(\sqrt{\pi}) + \frac{1}{2}
\end{aligned}$$

As $\sigma \to 0$ we have $H \to -\infty$

Sensitivity to the Choice of Units

$$H(N(0, \sigma)) = C + \ln \sigma$$

Differential entropy depends on the choice of units — a distributions on lengths will have a different entropy when measuring in inches than when measuring in feet.

Differential Cross-Entropy can Diverge to $-\infty$

Consider the unsupervised training objective.

$$\Phi^* = \underset{\Phi}{\operatorname{argmin}}\, E_{y\sim\text{train}} \, -\ln p_\Phi(y)$$

The training set is finite (discrete).

For each y the density $p_\Phi(y)$ can go to infinity.

This will drive the cross-entropy training loss to $-\infty$.

Differential Entropy Can Be Considered Infinite

An actual real number carries an infinite number of bits.

Consider quantizing the real numbers into bins.

A continuous probability densisty p assigns a probability $p(B)$ to each bin.

As the bin size decreases toward zero the entropy of the bin distribution increases toward ∞.

A meaningful convention is that $H(p) = +\infty$ for any continuous density p.

Differential KL-divergence is Meaningful

$$KL(p, q) = \int \left(\ln \frac{p(x)}{q(x)} \right) p(x) dx$$

Unlike differential entropy, differential KL divergence is always non-negative (but can be infinite).

Note that $KL(p, p) = 0$ independent of $H(p)$.

Mutual Information

$I(x, y)$ is the reduction in the number of bits we need to name y as a result of observing x (on average).

$$I(x, y) = E \ln \frac{P(x, y)}{P(x)P(y)}$$

$$= E \ln \frac{P(x, y)}{P(x)} - \ln P(y)$$

$$= H(y) - H(y|x)$$

Intuitively, how much does x know about y?

Differential Mutual Information

$$I(x, y) = KL(p(x, y), p(x)p(y))$$

$$= E_{x,y} \ln \frac{p(x, y)}{p(x)p(y)}$$

Mutual information is a KL divergence and hence differential mutual information is always non-negative.

The Data Processing Inequality

For continuous y and z with $z = f(y)$ we get that $H(z)$ can be either larger or smaller than $H(y)$ (consider $z = ay$ for $a > 1$ vs. $a < 1$).

However, mutual information is a KL divergence and is more meaningful than entropy and for $z = f(y)$ we do have

$$I(x, z) \leq I(x, y)$$

Continuous Cross-Entropy as Distortion

For a Gaussian VAE on images we typically have

$$\Phi^*, \Psi^* = \underset{\Phi,\Psi}{\mathrm{argmin}}\; E_{y\sim\mathrm{pop},z\sim\hat{p}_\Psi(z|y)}\; \ln\frac{\hat{p}_\Psi(z|y)}{p_\Phi(z)} - \ln p_\Phi(y|z)$$

$$p_\Phi(y|z) \;\propto\; \frac{1}{Z}\; \exp(||y - \hat{y}_\Phi(z)||^2/2\sigma^2)$$

$$\Phi^*, \Psi^* = \underset{\Phi,\Psi}{\mathrm{argmin}}\; E_{y\sim\mathrm{pop},z\sim\hat{p}_\Psi(z|y)}\; \ln\frac{\hat{p}_\Psi(z|y)}{p_\Phi(z)} + \frac{1}{2\sigma^2}||y - \hat{y}_\Phi(z)||^2 + d\ln\sigma$$

The KL divergence term is not problematic. The problematic differential cross-entropy term can just be thought of as a weighted L_2 distortion.

Continuous Cross-Entropy as Distortion

$$\Phi^*, \Psi^* = \operatorname*{argmin}_{\Phi, \Psi} \; E_{y \sim \text{pop}, z \sim \hat{p}_\Psi(z|y)} \; \ln \frac{\hat{p}_\Psi(z|y)}{p_\Phi(z)} \; {\color{red}+ \lambda \operatorname{Dist}(y, \hat{y}_\Phi(z))}$$

Various choices for distortion are possible including L_2 and L_1 distortion measures.

$$\operatorname{Dist}(y, \hat{y}) = ||y - \hat{y}||^2 \qquad (L_2)$$

$$\text{or} \;\; \operatorname{Dist}(y, \hat{y}) = ||y - \hat{y}||_1 = \sum_i |y[i] - \hat{y}[i]| \qquad (L_1)$$

END